

Juri Buora

Security and AI operations · early career

juribuora@gmail.com · juribuora.com · linkedin.com/in/juri-buora · github.com/JuriBuora

I build AI systems that do real work, and the safety around them, and I write all of it down. I am looking for a junior role in security or AI operations. I work with AI pair-programming: I set the goals, the guardrails and the review.

1,364

commits in 13 weeks on my workstation project

34

documented failure shapes in my release gate

243

days of public learning notes and labs

231

commits on a live client website

SELECTED WORK

WebCheckup

External, non-invasive website check-ups for small businesses: HTTPS, browser protections, mobile, speed, links. An email may state a finding only if the evidence reproduces. Italian and English, kept in sync by a test.

Foresight release gate

A pre-mortem and release check built after an AI system's defects reached real people. Each risky change names what could break, the mechanism that stops it, and a proof that was actually run.

Supervised agents, measured honestly

Approvals, audit records and crash recovery for an AI agent. In a 17-day live trial, 41 tasks ran and 2 distinct real tasks completed, so I switched it off and fixed the root causes.

Client work

Built and maintain a production website for a farm business. For a small business, analysed two legacy databases on verified copies and delivered a costed migration report and a first working workbook.

Fail-safe AI messaging

An auto-reply assistant with a filter at the delivery boundary. Measured with 123 scenarios run three times each; leaked reasoning, tool errors and echoed transcripts are blocked before sending.

Security review of my own tool

Found a DNS-rebinding path in a browser tool I wrote for an agent: the safety check and the browser could reach different addresses. Documented the options and their costs.

WHAT I BRING

Security

Web and network risk basics (TLS, headers, SSRF, DNS), least privilege, detection tuning practice in KQL, Linux and SQL labs, clear write-ups.

AI operations

Running agent gateways, routing between local and cloud models, output filtering, adversarial testing, evaluation with repeated runs.

Engineering

Git, React and TypeScript, Python, SQL, Docker, CI/CD on GitHub Actions, macOS and Linux. Italian and English.

HOW I WORK

- I look for what can go wrong before building, and put the control where it cannot be walked around.
- I verify what is actually running, not what the configuration says, and I report numbers that do not flatter me.
- I leave handoff notes so anyone can continue: what is done, what is left, and the next command.

Honest scope: this is solo work, done after hours, with AI pair-programming. Most repositories are private because they hold personal data; the numbers come from git and I will show them in a live walkthrough. Full detail, including projects I stopped and why: juribuora.com/what-im-doing